arXiv:2503.00260v2 [cs.CV] 24 Mar 2025

Seeing A 3D World in A Grain of Sand

Yufan Zhang[1] Yu Ji[2] Yu Guo[1] Jinwei Ye[1]
[1]George Mason University [2]LightThought LLC

https://miniature-3dgs.github.io/

Abstract

We present a snapshot imaging technique for recovering 3D surrounding views of miniature scenes. Due to their intricacy, miniature scenes with objects sized in millimeters are difficult to reconstruct, yet miniatures are common in life and their 3D digitalization is desirable. We design a catadioptric imaging system with a single camera and eight pairs of planar mirrors for snapshot 3D reconstruction from a dollhouse perspective. We place paired mirrors on nested pyramid surfaces for capturing surrounding multi-view images in a single shot. Our mirror design is customizable based on the size of the scene for optimized view coverage. We use the 3D Gaussian Splatting (3DGS) representation for scene reconstruction and novel view synthesis. We overcome the challenge posed by our sparse view input by integrating visual hull-derived depth constraint. Our method demonstrates state-of-the-art performance on a variety of synthetic and real miniature scenes.

1. Introduction

To see a World in a Grain of Sand
And a Heaven in a Wild Flower,
Hold Infinity in the palm of your hand
And Eternity in an hour.

- William Blake

Most existing works on 3D reconstruction or novel view synthesis focus on large-scale or life-size scenes. The tiny world of miniatures (*e.g.*, objects sized in centimeters or even millimeters) is somewhat neglected. 3D reconstruction of miniature scenes is challenging due to limitations on lenses, image resolution and reconstruction accuracy. To take images of miniatures, one would need a macro lens to magnify the tiny objects, such that they could cover substantial amount of pixels in an image. Due to high magnification ratio, macro images usually have shallow depth of field, making it hard to capture all-in-focus image. In addition, many miniatures have little textures (as limited by their size), which poses challenge to photogrammetry-based 3D reconstruction. Yet miniatures scenes are common in life (see examples in Fig. 1): from toys and decorations in household to artisanal crafts and antiques in design studios and museums, high-quality 3D reconstruction not only opens up new ways for viewing and interacting with those miniatures, and also benefits their preservation [4] .

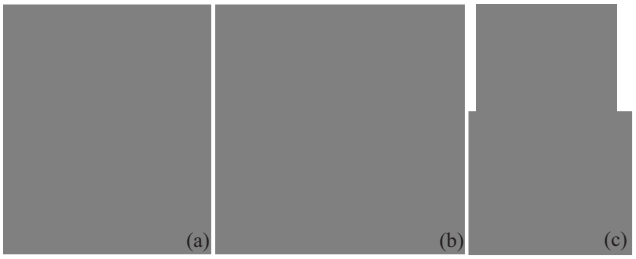


Figure 1. (a) A miniature scene with tiny object; (b) Example image captured by our catadioptric imaging system; (c) Novel view synthesis results.

In this work, we present a snapshot solution for reconstructing miniature scenes and synthesizing their images in full surround 360° views. We design a novel planar catadioptric lens to enable zoom-in all-focus imaging and time-synchronized full surround multi-view acquisition. Specifically, we use the reflection between paired mirrors to fold the light path in order to zoom-in onto small object in short distance. We arrange eight pairs of mirrors on the sides of two nested octagonal pyramids to capture surrounding views of the scene in a single shot. We analyze the ray geometry of our lens and derive closed-form formulas for optimizing the mirror configuration based on the size of the scene, such that the multi-view coverage of the scene is optimal. We map our multi-view sub-images to virtual cameras and pre-calibrate the virtual camera parameters. The calibration only needs to run once when the mirror configuration is set. Since our calibrated camera parameters are highly accurate, they greatly benefit scene reconstruction. In contrast, self-calibration methods (*e.g.*, structure from motion [27]) heavily rely on scene features and are not reliable for miniature scenes with little textures.

We use the 3D Gaussian Splatting (3DGS) [12] for scene

representation, and synthesize full surround novel views using our multi-view images and pre-calibrated camera parameters. Since our input views are sparse (*e.g.*, 8 views on a 360° circle), we use additional depth constraints derived from visual hull to improve the 3DGS reconstruction. Specifically, we first calculate visual hull using the object's silhouette masks extracted from multi-view images. We then project the visual hull to depth map under each viewpoint. Since visual hull is a convex volume that fully enclose the object, we propose a weighted depth loss that penalizes more on depth values greater than the visual hull depth (which indicates the point is outside of the visual hull). Our depth loss is especially effective for surfaces with little textures, and it can be used for generic 3DGS when visual hulls are available.

We validate our method on both synthetic and real data. For synthetic experiments, we render images by emulating our imaging setup and quantitatively evaluate the novel view synthesis results. For real experiments, we custom-build lens prototypes using 3D-printed housing and acrylic mirrors, and capture images of a variety of miniature scenes. We compare our method with recent sparse-view 3DGS approaches, and demonstrate better rendering quality. Our main contributions are summarized as follows:

- We design a planar catadioptric imaging system for recovering 3D surrounding views of miniature scenes in a single shot.
- We analyze the ray geometry of our catadioptric lens and derive closed-form formulas for calculating optimized lens parameters given scene information.
- We propose a novel weighted depth loss based on visual hull to improve 3DGS with sparse view input.
- We build a prototype for the proposed imaging system and validate our approach on real miniature scenes.

2. Related Work

Here we briefly review prior works that are most relevant to our imaging system and reconstruction algorithm.

Mirror-based imaging system. Mirrors, either planar or curved, are widely used for building imaging systems with extended field-of-view or snapshot multi-view capacities. Imaging systems that involve both mirrors and refractive lenses are called catadioptric systems. They have been extensively studied and used for stereo [23, 28, 35], panoramic [3, 29], surround-view [1, 17] and light field [6, 10, 22] imaging. Reshetouski and Ihrke [25] provide a comprehensive survey on the design and applications of various mirror-based imaging systems. Gluckman and Nayar [7, 8] show all possible configurations for catadioptric stereo and derive the reflection transformation for mirror-based image formation. Notably, kaleidoscopic imaging systems use the inter-reflection of mirrors to generate variations in viewpoints and illumination for 3D shape [26, 33] and reflectance [9, 11] reconstruction. Various techniques are explored for 3D reconstruction, including space carving [26], multi-view stereo [20], structured light [1] and neural surface representation [2]. Our work uses eight pairs of mirrors for full-surround imaging and the 3DGS framework for scene reconstruction. In contrast to kaleidoscopes, we avoid inter-reflection in our mirror system. Thus, our multi-view images are much easier to separate and analyze, although our angular resolution is sacrificed.

3DGS with sparse view input. 3D Gaussian Splatting (3DGS) [12] is an efficient radiance field representation that allows for high-quality real-time rendering of novel views, with relatively short training time. But the original 3DGS requires dense view supervision to achieve high quality rendering. With few reference views, the algorithm tends to overfit on inputs, resulting artifacts in unseen views. The problem gets even worse for surround-view rendering. Many recent techniques are proposed to allow robust 3DGS-based novel view synthesis when few views are available. A major trend is to use depth maps to provide additional supervision. FSGS [42] and SparseGS [37] use monocular depth predicted by pre-trained model [24]. GS2Mesh [34] and InstantSplat [5] adopt depth learned from multi-view stereo. ReconFusion [36] and GaussianObject [38] use diffusion model to synthesize addition views in order to provide dense viewpoint and depth supervision. RaDe-GS [40] improves the depth rasterization in 3DGS. In this work, we propose a new weighted depth loss that leverages depth map generated from visual hull to improve the quality of novel view synthesis with sparse surround-view input.

3. Imaging System

In this section, we introduce our imaging system for snapshot full-surround miniature scene reconstruction. We first describe the design of our planar catadioptric lens, and then analyze the ray geometry of the lens and derive optimized lens parameters given scene information.

3.1. Catadioptric Lens Design

Our imaging system consists of a camera and a planar catadioptric lens that allows for synchronized full-surround multi-view image acquisition. The conceptual design is illustrated in Fig. 2. Our catadioptric lens is made with eight pairs of flat mirrors that are circularly arranged on the sides of two nested octagonal pyramids. In each pair, the mirrors are facing each other with different tilting angles and shifted locations. Through two times of reflections, the mirror pair guides light from a scene that is underneath to the viewing camera on top. In this way, the light paths are folded

Appendices

A. Geometric Derivations

A.1. Derivation of Angles in Section 3.2

Here we show how to derive the angles we used in Section 3.2, when formulating the effective viewing volume. Our goal is derive the half apex angle of the viewing volume (θ), given the tilting angles of the two mirrors (α_1 and α_2). For ease of reference, we introduce auxiliary angles labeled in numbers. All the angles that we have referred to are annotated in Fig. 11.

Since $\angle 1$ and $\angle 2$ are vertical angles, we have $\angle 2 = \angle 1 = 90° - \alpha_1$. Since $\angle 2$ and ω_1 are complementary, we can calculate the incident/exit angle of reflection on M_1 as:

$$\omega_1 = 90° - \angle 2 = \alpha_1. \quad (8)$$

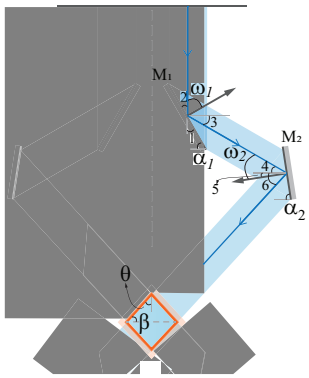


Figure 11. Angle annotations.

Since $\angle 3$ and $\angle 4$ are alternate angles, we have $\angle 4 = \angle 3 = 2\omega_1 - 90°$. By substituting ω_1 with Eq. 8, we have $\angle 4 = 2\alpha_1 - 90°$. Since $\angle 5$ and α_2 are complementary, we have $\angle 5 = 90° - \alpha_2$. Therefore, the incident/exit angle of reflection on M_2 can be calculated as:

$$\omega_2 = \angle 4 + \angle 5 = 2\alpha_1 - \alpha_2. \quad (9)$$

Since $\angle 6$ and β are congruent, we have $\beta = \angle 6 = \angle 5 + \omega_2$. By substituting ω_2 and $\angle 5$, we have $\beta = 90° - 2\Delta\alpha$, where $\Delta\alpha = \alpha_2 - \alpha_1$. The half apex angle of the effective viewing volume, being complenentary to β, is thus:

$$\theta = 90° - \beta = 2\Delta\alpha. \quad (10)$$

A.2. Derivation of Conditions in Section 3.2

Derivation of condition (i). This condition is introduced to allow light to travel through the lens from one end to the other, after being reflected by the two mirrors in a pair. In addition, the multi-view images formed by the eight mirror pairs should have overlaps, in order to be practical for scene reconstruction.

With $\alpha_1 > 45°$ and $\alpha_2 < 90°$, we guarantee that light from the scene could travel through our mirror lens and reach the camera on the other end (*i.e.*, light path wouldn't turn around inside of the lens). If $\alpha_2 < \alpha_1$, the light exiting the lens (after reflected by M_2) would be diverging (see Fig. 12 (a)), resulting none-overlapping multi-view images. So we have $45° < \alpha_1 < \alpha_2 < 90°$.

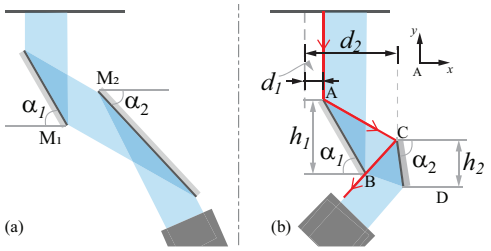


Figure 12. (a) The diverging situation when $\alpha_2 < \alpha_1$; (b) The extreme situation without inter-reflection, in which the reflected ray from M_2 intersects with the bottom edge of M_1.

Derivation of condition (ii). Here we derive the minimum vertically projected height of M_2 (denoted as h_2), such that it can cover the entire light beam reflected from M_1.

The width of parallel light beam reflected from M_1 is $w = h_1 / \tan\alpha_1$, where h_1 is the vertically projected height of M_1. In order to cover the entire beam, the length of M_2 (denoted as l_2) should satisfy:

$$l_2 \geq \frac{w}{\sin(\alpha_2 - \angle 4)} = \frac{w}{\cos(2\alpha_1 - \alpha_2)}. \quad (11)$$

Substituting $l_2 = h_2 / \sin\alpha_2$ and $w = h_1 / \tan\alpha_1$, we can rewrite Eq. 11 as:

$$h_2 \geq \frac{\sin\alpha_2}{\tan\alpha_1 \cdot \cos(2\alpha_1 - \alpha_2)} \cdot h_1. \quad (12)$$

Derviation of condition (iii). Here we derive the minimum separation between the two mirrors in order to avoid inter-reflection. We quantify this distance as $d_2 - d_1$ (where d_1 and d_2 are the distances from M_1 and M_2's upper edges to the central ray), when given their vertically projected heights h_1, h_2 and tilting angles α_1, α_2. We consider the extreme situation when the leftmost ray of the light beam intersects with the bottom edge of M_1 after reflecting from M_2 (see Fig. 12 (b)).

We denote the end points of M_1 and M_2 in the 2D cross-section plot as A, B, C, and D. We setup a coordinate system with A as the origin as shown in Figure 12(b). The line equation for M_1 (line AB) can be written as:

$$y = -\tan\alpha_1 \cdot x. \quad (13)$$

The line equation for the leftmost ray incident to M_2 (line AC) can be written as:

$$y = \cot 2\alpha_1 \cdot x. \quad (14)$$

Since $x_C = d_2 - d_1$, we plug it into Eq. 14 and calculate the coordinate of C as $(d_2 - d_1, \cot 2\alpha_1 \cdot (d_2 - d_1))$. The line equation for leftmost ray reflected from M_2 (line BC) can thus be calculated as:

$$y = \cot\Delta\alpha \cdot (x - (d_2 - d_1)) + \cot 2\alpha_1 \cdot (d_2 - d_1), \quad (15)$$

where $\Delta\alpha = \alpha_2 - \alpha_1$. By combining Eq. 13 and Eq. 15, we can calculate the x coordinate of B as:

$$x_B = \frac{\cot 2\Delta\alpha - \cot 2\alpha_1}{\tan\alpha_1 + \cot 2\Delta\alpha} \cdot (d_2 - d_1). \quad (16)$$

To avoid inter-reflection, x_B should be satisfy: $x_B \geq h_1/\tan\alpha_1$. Subsituting x_B with Eq. 16, we obtain the third condition regarding the mirror distances:

$$d_2 \geq \frac{\tan\alpha_1 + \cot 2\Delta\alpha}{\tan\alpha_1 \cdot (\cot 2\Delta\alpha - \cot 2\alpha_1)} \cdot h_1 + d_1. \quad (17)$$

B. More Details on Pre-processing Steps

Fig. 13 shows how our captured raw image is processed into multi-view input to 3DGS. A raw image captured by our portable lens prototype is shown in Fig 13 (a). Its resolution is 2448×2048. We first apply a multi-view mask to extract the effective regions formed through mirror reflection. The filtered image is shown in Fig. 13(b). Then, for each sub-image, we re-project it to allow smooth view transition (we update camera poses after re-projection). We also mask out the background and only reconstruct the foreground objects. The processed image for one sub-view (highlighted in red) is shown in Fig. 13 (c). This image is with resolution 800×800. The eight sub-view images processed in this way are used as input to 3DGS.

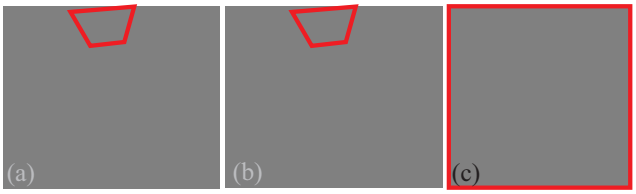


Figure 13. (a) Our captured raw image; (b) Image filtered by the multi-view mask; (c) Re-projected image of the highlighted view.

C. Lens Design Comparison

Here we show comparison between two lens designs with different mirror configuration. Prototypes of the two designs are shown in Fig. 14. The two lenses have the same base lengths for the inner and outer pyramids, with different tilting angles for the mirrors. The parameters we use are shown in Table 2. Images taken with the two lenses are shown in Fig. 14.

We can see that design (b), which has larger $\Delta\alpha$, has better coverage of the side views (*e.g.*, the figurine's face becomes visible in (b)). This is equivalent to having virtual cameras with more oblique angles. Such configuration is preferred since it provides fuller coverage of the scene. This observation is consistent with our guidelines on optimizing the mirror configuration.



Figure 14. Comparison between two lens designs. Here we show the lens prototypes and their captured images with zoom-in views.

Table 2. Mirror parameters of the two different designs.

	α_1	α_2	$\Delta\alpha$
Design (a)	75°	85°	10°
Design (b)	60°	85°	25°

D. Additional Experimental Results

D.1. Ablation on Depth Loss

Fig. 15 compares depth maps obtained by different methods for a real scene (*i.e.*, the "frog" scene). Specifically, the MiDaS [24] depth is used by FSGS [42]; Depth Anything V2 [39] is used by Hierarchical 3DGS [13]; and the visual hull depth is used by our approach. We can see that Depth Anything provides much better depth prior than MiDaS depth. Our visual hull depth outperforms Depth Anything result in details (*e.g.*, the frog legs have more discernible depth variation in the visual hull depth). Moreover, the visual hull projection provides depth values in absolute scale, whereas the other two learning-based methods estimate relative depths.

We performed an ablation study on depth loss using the "skull" scene (see Fig. 16). We compare the PSNR of synthesized novel views for three variants of our algorithm: without depth loss, with monocular depth (Depth Anything V2 [39] depth), and with visual hull depth (VH depth). The table below shows the ablation study on depth loss.

Table 3. Ablation study on depth loss.

Variant	w/o depth	w. monodepth	w. VH depth
PSNR	29.2052	29.3698	29.3856

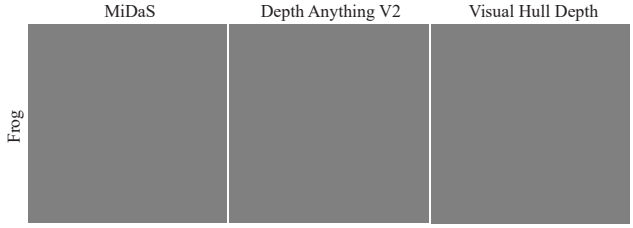


Figure 15. Comparison of depth map obtained by different methods.

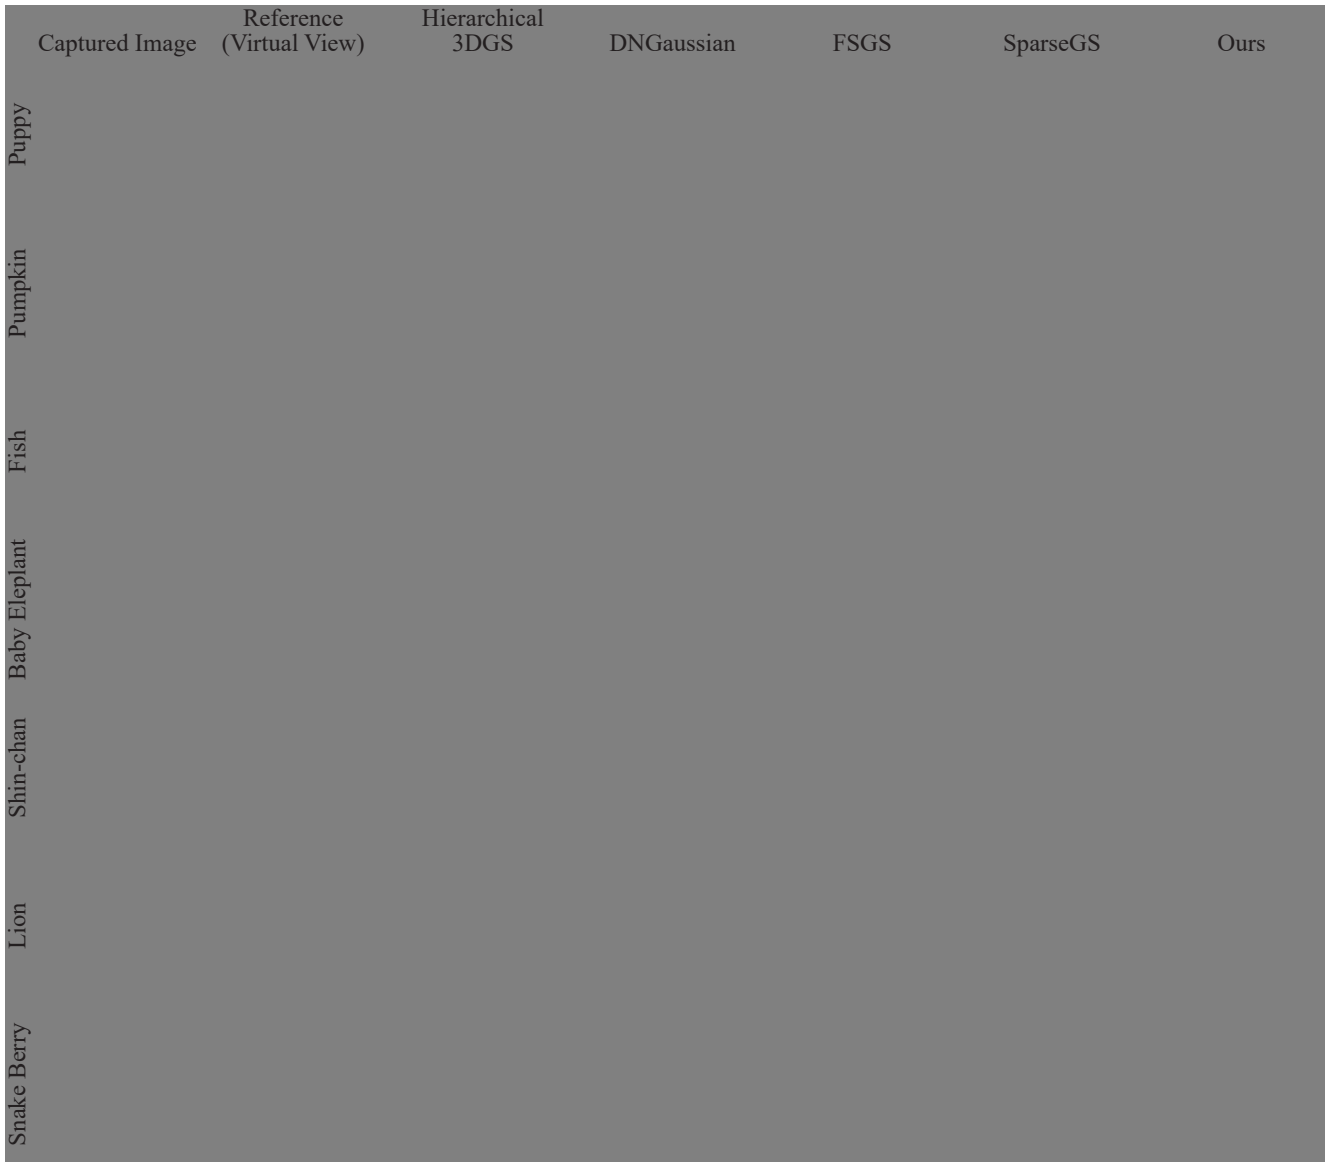


Figure 17. Additional visual comparison results on real data.

References

[1] Byeongjoo Ahn, Ioannis Gkioulekas, and Aswin C Sankaranarayanan. Kaleidoscopic structured light. *ACM TOG*, 40 (6):1–15, 2021. 2, 3

[2] Byeongjoo Ahn, Michael De Zeeuw, Ioannis Gkioulekas, and Aswin C. Sankaranarayanan. Neural kaleidoscopic space sculpting. In *CVPR*, 2023. 2

[3] Javaan Singh Chahl and Mandyam Veerambudi Srinivasan. Reflective surfaces for panoramic imaging. *Applied optics*, 36(31):8275–8285, 1997. 2

[4] Jack Davy and Charlotte Dixon. *Worlds in Miniature*. UCL Press, 2019. 1

[5] Zhiwen Fan, Wenyan Cong, Kairun Wen, Kevin Wang, Jian Zhang, Xinghao Ding, Danfei Xu, Boris Ivanovic, Marco Pavone, Georgios Pavlakos, et al. Instantsplat: Unbounded sparse-view pose-free gaussian splatting in 40 seconds. *arXiv preprint arXiv:2403.20309*, 2, 2024. 2, 8

[6] Martin Fuchs, Markus Kächele, and Szymon Rusinkiewicz. Design and fabrication of faceted mirror arrays for light field capture. *Comput. Graph. Forum*, 32, 2013. 2

[7] Joshua Gluckman and Shree Nayar. Rectified catadioptric stereo sensors. In *CVPR*, 2000. 2

[8] Joshua Gluckman and Shree Nayar. Catadioptric stereo using planar mirrors. *IJCV*, 44:65–79, 2001. 2

[9] Jefferson Y. Han and Ken Perlin. Measuring bidirectional texture reflectance with a kaleidoscope. *ACM TOG*, 22(3), 2003. 2

[10] Ivo Ihrke, Timo Stich, Heiko Gottschlich, Marcus Magnor,

and Hans-Peter Seidel. Fast incident light field acquisition and rendering. *Journal of WSCG*, 16:25–32, 2008. 2

[11] Ivo Ihrke, Ilya Reshetouski, Alkhazur Manakov, Art Tevs, Michael Wand, and Hans-Peter Seidel. A kaleidoscopic approach to surround geometry and reflectance acquisition. In *CVPRW*, 2012. 2, 3

[12] Bernhard Kerbl, Georgios Kopanas, Thomas Leimkühler, and George Drettakis. 3D gaussian splatting for real-time radiance field rendering. *ACM TOG*, 42(4), 2023. 1, 2, 5

[13] Bernhard Kerbl, Andreas Meuleman, Georgios Kopanas, Michael Wimmer, Alexandre Lanvin, and George Drettakis. A hierarchical 3d gaussian representation for real-time rendering of very large datasets. *ACM TOG*, 43(4), 2024. 6, 7, 10, 11

[14] Jungho Kim, Kuk-jin Yoon, Jun-sik Kim, and Inso Kweon. Visual slam by single-camera catadioptric stereo. In *SICE-ICASE International Joint Conference*, 2006. 3

[15] Alexander Kirillov, Eric Mintun, Nikhila Ravi, Hanzi Mao, Chloe Rolland, Laura Gustafson, Tete Xiao, Spencer Whitehead, Alexander C. Berg, Wan-Yen Lo, Piotr Dollár, and Ross Girshick. Segment anything. *arXiv:2304.02643*, 2023. 5

[16] Muhammed Kocabas, Jen-Hao Rick Chang, James Gabriel, Oncel Tuzel, and Anurag Ranjan. Hugs: Human gaussian splats. In *CVPR*, pages 505–515, 2024. 5

[17] Douglas Lanman, Daniel Crispell, and Gabriel Taubin. Surround structured lighting for full object scanning. In *International Conference on 3-D Digital Imaging and Modeling (3DIM)*, 2007. 2

[18] Aldo Laurentini. The visual hull concept for silhouette-based image understanding. *IEEE TPAMI*, 16(2):150–162, 1994. 5

[19] Jiahe Li, Jiawei Zhang, Xiao Bai, Jin Zheng, Xin Ning, Jun Zhou, and Lin Gu. Dngaussian: Optimizing sparse-view 3d gaussian radiance fields with global-local depth normalization. In *CVPR*, 2024. 7, 11

[20] Adrien Mas, Guillaume Druart, Maxime Vaché, Sylvain Favier, Aymeric Alazarine, Eric Compain, Nathalie Morin, and Nicolas Guérineau. Kaleidoscope-based multi-view infrared system. *Optics Letters*, 44(20):4977–4980, 2019. 2

[21] Ben Mildenhall, Pratul P. Srinivasan, Matthew Tancik, Jonathan T. Barron, Ravi Ramamoorthi, and Ren Ng. Nerf: Representing scenes as neural radiance fields for view synthesis. In *ECCV*, 2020. 5

[22] Yasuhiro Mukaigawa, Seiichi Tagawa, Jaewon Kim, Ramesh Raskar, Yasuyuki Matsushita, and Yasushi Yagi. Hemispherical confocal imaging using turtleback reflector. In *ACCV*, 2011. 2

[23] Sameer A. Nene and Shree K. Nayar. Stereo with mirrors. *ICCV*, pages 1087–1094, 1998. 2

[24] René Ranftl, Katrin Lasinger, David Hafner, Konrad Schindler, and Vladlen Koltun. Towards robust monocular depth estimation: Mixing datasets for zero-shot cross-dataset transfer. *IEEE TPAMI*, 44(3), 2022. 2, 10

[25] Ilya Reshetouski and Ivo Ihrke. *Mirrors in computer graphics, computer vision and time-of-flight imaging*. 2013. 2

[26] Ilya Reshetouski, Alkhazur Manakov, Hans-Peter Seidel, and Ivo Ihrke. Three-dimensional kaleidoscopic imaging. In *CVPR*, 2011. 2

[27] Johannes Lutz Schönberger and Jan-Michael Frahm. Structure-from-motion revisited. In *CVPR*, 2016. 1, 5

[28] Tomáš Svoboda and Tomáš Pajdla. Epipolar geometry for central catadioptric cameras. *IJCV*, 49:23–37, 2002. 2

[29] Rahul Swaminathan, Shree K Nayar, and Michael D Grossberg. Framework for designing catadioptric projection and imaging systems. *ICCV*, 2003. 2

[30] Jiaxiang Tang, Jiawei Ren, Hang Zhou, Ziwei Liu, and Gang Zeng. Dreamgaussian: Generative gaussian splatting for efficient 3d content creation. *ICLR*, 2024. 5

[31] Eric J. Tremblay, Ronald A. Stack, Rick L. Morrison, and Joseph E. Ford. Ultrathin cameras using annular folded optics. *Appl. Opt.*, 46(4):463–471, 2007. 3

[32] Eric J. Tremblay, Ronald A. Stack, Rick L. Morrison, Jason Harris Karp, and Joseph E. Ford. Ultrathin four-reflection imager. *Applied optics*, 48 2:343–54, 2009. 3

[33] Felix Wechsler, Rainer Heintzmann, and Ivo Ihrke. Kaleidomicroscope-a kaleidoscopic multiview microscope. In *Computational Optical Sensing and Imaging*, pages CTu4F–5, 2022. 2

[34] Yaniv Wolf, Amit Bracha, and Ron Kimmel. Gs2mesh: Surface reconstruction from gaussian splatting via novel stereo views. In *ECCV*, 2024. 2

[35] Hsien-Huang P. Wu and Shih-Hsin Chang. Design of stereoscopic viewing system based on a compact mirror and dual monitor. *Optical Engineering*, 49:027401, 2010. 2

[36] Rundi Wu, Ben Mildenhall, Philipp Henzler, Keunhong Park, Ruiqi Gao, Daniel Watson, Pratul P Srinivasan, Dor Verbin, Jonathan T Barron, Ben Poole, and Aleksander Holynski. Reconfusion: 3d reconstruction with diffusion priors. In *CVPR*, 2024. 2

[37] Haolin Xiong, Sairisheek Muttukuru, Rishi Upadhyay, Pradyumna Chari, and Achuta Kadambi. SparseGS: Real-time 360° sparse view synthesis using gaussian splatting. *Arxiv*, 2023. 2, 7, 11

[38] Chen Yang, Sikuang Li, Jiemin Fang, Ruofan Liang, Lingxi Xie, Xiaopeng Zhang, Wei Shen, and Qi Tian. Gaussianobject: High-quality 3d object reconstruction from four views with gaussian splatting. *ACM TOG*, 43(6), 2024. 2

[39] Lihe Yang, Bingyi Kang, Zilong Huang, Zhen Zhao, Xiaogang Xu, Jiashi Feng, and Hengshuang Zhao. Depth anything v2. *arXiv:2406.09414*, 2024. 10

[40] Baowen Zhang, Chuan Fang, Rakesh Shrestha, Yixun Liang, Xiaoxiao Long, and Ping Tan. Rade-gs: Rasterizing depth in gaussian splatting. *arXiv preprint arXiv:2406.01467*, 2024. 2

[41] Xiaoyu Zhou, Zhiwei Lin, Xiaojun Shan, Yongtao Wang, Deqing Sun, and Ming-Hsuan Yang. Drivinggaussian: Composite gaussian splatting for surrounding dynamic autonomous driving scenes. In *CVPR*, pages 21634–21643, 2024. 5

[42] Zehao Zhu, Zhiwen Fan, Yifan Jiang, and Zhangyang Wang. FSGS: Real-time few-shot view synthesis using gaussian splatting. In *ECCV*, 2024. 2, 7, 10, 11